

MEMORY-IMPOSTOR ROLLOUTS IN RETRIEVAL-AUGMENTED WORLD MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

Memory-augmented world models retrieve episodic transitions, trajectories, or cases to improve prediction under sparse data. This paper studies a narrow failure at the interface of retrieval and rollout selection: stale or support-mismatched memory can create *memory-impostor* rollouts, candidates whose predicted value is high because retrieved cases come from the wrong latent regime while their true executed value is low. Increasing the candidate budget then makes the planner more likely to select those impostors. We give a finite- N mechanism statement, propose provenance diagnostics based on retrieval precision, staleness, counterfactual memory dropout, hallucinated-rollout rate, and proxy-true gap, and instantiate the mechanism in a controlled latent-regime world. In the main stale-memory condition, naive memory selection improves its proxy score while true return falls from -7.87 at $N = 1$ to -13.33 at $N = 64$; selected rollouts have 15.7% retrieval precision, 84.3% stale retrieval, and 5.6% hallucinated-rollout rate. A repair using recency-aware retrieval, disagreement-aware memory gating, and dropout-risk penalties reaches -1.90 true return at the same N , narrowing but not eliminating the gap to an oracle retriever. The claim is intentionally scoped: this is a runnable diagnostic for an architecture-specific failure mode, not evidence that memory is generally harmful.

1 INTRODUCTION

World models let agents reason in imagined futures rather than only through environment interaction (Sutton, 1990; Ha & Schmidhuber, 2018; Hafner et al., 2019; 2020; Schrittwieser et al., 2020). A common inference pattern is to sample many candidate rollouts and select the one with the highest model score. This is attractive whenever the generator is cheap and the scoring proxy is believed to be reliable.

External memory gives a second source of test-time power. Episodic control, retrieval-augmented reinforcement learning, case-based reasoning, kNN language models, and retrieval-augmented generation all show that non-parametric memory can improve adaptation and factuality (Aamodt & Plaza, 1994; Pritzel et al., 2017; Wayne et al., 2018; Khandelwal et al., 2020; Guu et al., 2020; Lewis et al., 2020; Goyal et al., 2022). For world models, the hope is direct: retrieve relevant cases and make better local transition predictions.

This paper asks what happens when these two useful ideas are composed under support mismatch. If the retriever returns stale cases from a latent regime with the wrong transition sign, a memory-augmented model can assign high value to a candidate future that the real environment will not realize. Max-score rollout selection then searches for that error. More samples need not make the selected rollout more grounded; they can make it a better impostor under the memory-biased proxy.

Our contributions are:

- we define the memory-impostor rollout failure mode for retrieval-augmented world-model planning;
- we give a finite- N mechanism statement showing why larger candidate budgets amplify high-proxy, low-true memory impostors;

- we introduce provenance diagnostics for selected-rollout retrieval precision, staleness, memory dropout variance, hallucinated-rollout rate, and memory exploitation gap;
- we instantiate the mechanism in a controlled synthetic world where latent dynamics, memory provenance, and true rollout utility are all known;
- we test a simple non-oracle repair combining recency-aware retrieval, disagreement-aware memory gating, and an uncertainty penalty.

2 RELATED WORK

World models and model-based RL. Model-based RL has long used learned dynamics for planning (Sutton, 1990). Modern world models learn latent dynamics from high-dimensional observations and use imagined trajectories for control (Ha & Schmidhuber, 2018; Hafner et al., 2019; 2020). Probabilistic and offline model-based methods such as PETS, MBPO, MOREL, and MOPO explicitly address model bias and uncertainty (Chua et al., 2018; Janner et al., 2019; Kidambi et al., 2020; Yu et al., 2020). Our setting is not a new planner; it isolates a failure source that arises when the local dynamics model is partly retrieved from an external case memory.

Episodic, retrieval, and case-based memory. Case-based reasoning formalized retrieve/reuse/revise/retain cycles (Aamodt & Plaza, 1994). Neural episodic control and MERLIN showed that fast memory can help agents adapt (Pritzel et al., 2017; Wayne et al., 2018), while retrieval-augmented RL and model-based episodic memory bring retrieved experience into control and model-based settings (Goyal et al., 2022; Le et al., 2021). Retrieval-augmented language models and kNN-LMs show a closely related non-parametric interpolation pattern (Khandelwal et al., 2020; Guu et al., 2020; Lewis et al., 2020; Borgeaud et al., 2022; Wu et al., 2022). These works motivate memory augmentation; our contribution is to show how inference-time selection can amplify a memory error that average prediction metrics may understate.

Test-time selection and proxy optimization. Self-consistency, verifier ranking, Tree of Thoughts, and test-time compute scaling all improve outputs by sampling or searching multiple candidates (Wang et al., 2023; Cobbe et al., 2021; Yao et al., 2023; Snell et al., 2024). Reward-model overoptimization shows that selecting against an imperfect proxy can degrade true utility (Gao et al., 2023). Our finite-N argument is simple in that tradition; the novelty is the memory-specific source of proxy/true inversion and the corresponding retrieval diagnostics.

3 MECHANISM

Let a candidate rollout be grounded with probability $1 - p$ and a memory-impostor with probability $p > 0$. A memory-impostor is a rollout whose model score is high because retrieved cases are stale or support mismatched, but whose true return is low under the environment dynamics. Suppose every impostor has proxy score s_I and true utility u_I , and every grounded candidate has proxy score s_G and true utility u_G , with $s_I > s_G$ but $u_I < u_G$.

Proposition 1 (Finite-N memory-impostor amplification). *If a max-score selector chooses the candidate with highest proxy score from N independent samples under the two-type model above, then*

$$P(\text{impostor selected}) = 1 - (1 - p)^N,$$

and the expected true utility of the selected candidate is

$$\mathbb{E}[u_N] = (1 - (1 - p)^N) u_I + (1 - p)^N u_G,$$

which decreases monotonically in N .

The proposition is not presented as deep probability theory. Its purpose is to turn a reviewer’s likely objection into a diagnostic: if memory retrieval can produce high-proxy low-true candidates, increasing N should increase proxy score, selected memory staleness, and proxy-true gap while decreasing true return.




Figure 1: In the stale-memory condition, naive memory selection improves model proxy score as N grows while true return degrades. The repaired selector tracks the non-memory/oracle direction because it suppresses stale retrieved transitions.

4 DIAGNOSTIC AND REPAIR

For a selected rollout τ , define the memory exploitation gap

$$\text{MEG}(\tau) = S_{\text{model}}(\tau) - U_{\text{true}}(\tau),$$

where S_{model} is the model proxy used for selection and U_{true} is the environment return used only for evaluation. We also log retrieval precision, stale retrieval rate, disagreement among retrieved transition signs, and counterfactual memory dropout variance: the variance of predicted deltas under leave-one-retrieved-case-out perturbations.

The repair intentionally uses only observable memory metadata and transition statistics, not the hidden regime label. It has three parts. First, retrieval is recency-aware, so recent calibration transitions can compete with dense stale cases. Second, memory interpolation is gated down when retrieved cases disagree or are old. Third, the selection proxy is penalized by the average dropout-risk score. This is a small repair rather than a claim of optimality; the point is to test whether the diagnosis points toward an effective intervention.

5 CONTROLLED BENCHMARK

The benchmark is a one-dimensional latent-regime world. The observed state is position x and actions are $a \in \{-1, 0, 1\}$. The current hidden regime is *reverse*, so the true transition is $x_{t+1} = x_t - a_t$. The memory store is dense with stale *forward* transitions where $x_{t+1} = x_t + a_t$ and sparse with recent reverse calibration transitions. A naive nearest-neighbor retriever sees only (x, a) , so the stale forward cases often dominate.

Candidate action sequences are sampled with both directional modes. Each model rolls out every candidate, scores the predicted final distance to a fixed goal, and executes the proxy-best sequence under the true reverse dynamics. We compare four selectors: a non-memory model fit from recent calibration transitions, naive memory, oracle memory that is allowed to filter by the hidden regime, and the repaired memory model.

6 RESULTS

Figure 1 shows the main scaling pattern. At stale fraction 0.85, naive memory selection falls from -7.87 true return at $N = 1$ to -13.33 at $N = 64$, even though its selected proxy score is -2.59. The resulting exploitation gap is 10.74. The non-memory baseline reaches -0.18, the repaired model reaches -1.90, and the oracle retriever reaches -1.73.

Figure 2 confirms that the failure is not merely bad random actions. At $N = 64$, naive selected rollouts have 15.7% retrieval precision and 84.3% stale retrieval. Their hallucinated rollout rate is




Figure 2: Selected-rollout diagnostics reveal the mechanism. Naive memory selections have poor retrieval precision and a high hallucinated-rollout rate at large N ; the repair improves selected retrieval quality and sharply reduces hallucinated rollouts.




Figure 3: At the largest N , the repair improves true return by 11.44 over naive memory. The remaining oracle gap is 0.17, which is expected because the oracle can directly filter by hidden regime.

5.6%. The repaired model improves selected retrieval precision to 50.9%, lowers stale retrieval to 49.1%, and reduces hallucinated rollouts to 0.0%.

Figure 3 makes the repair boundary visible. The repaired model is not an oracle and does not claim to recover the hidden regime perfectly. It uses age, disagreement, and memory-dropout risk as support diagnostics, which are available to a memory system without revealing the environment’s latent label.

7 LIMITATIONS

The benchmark is synthetic and deliberately low-dimensional. That is a strength for mechanism isolation and a weakness for external validity. The repair is also minimal: realistic systems would need stronger retrieval calibration, diversity constraints, causal or regime-aware memories, and task-specific uncertainty models. Finally, the formal claim is only the two-type finite- N memory-impostor amplification mechanism; it does not prove that a particular neural retriever will produce impostors in deployment.

8 REPRODUCIBILITY AND CHECKLIST

The repository includes the full synthetic generator, tests, raw result CSVs, plotting scripts, the 100+ entry related-work matrix, and this anonymous ICLR source. The oracle baseline is explicitly marked as using the hidden latent regime, while the repair is marked as non-oracle. No human-subject data or private datasets are used. Compute is a CPU-only synthetic run.

REFERENCES

- Agnar Aamodt and Enric Plaza. Case-based reasoning: Foundational issues, methodological variations, and system approaches. *AI Communications*, 1994.
- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego de Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, Andy Brock, Michela Paganini, Geoffrey Irving, Oriol Vinyals, Simon Osindero, Karen Simonyan, Jack W. Rae, Erich Elsen, and Laurent Sifre. Improving language models by retrieving from trillions of tokens. In *International Conference on Machine Learning*, 2022.
- Kurtland Chua, Roberto Calandra, Rowan McAllister, and Sergey Levine. Deep reinforcement learning in a handful of trials using probabilistic dynamics models. In *Advances in Neural Information Processing Systems*, 2018.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. *International Conference on Machine Learning*, 2023.
- Anirudh Goyal, Abram L. Friesen, Theophane Weber, Andrea Banino, Nan Rosemary Ke, Adria Puigdomenech Badia, Arthur Guez, Mehdi Mirza, Peter C. Humphreys, Ksenia Konyushkova, Laurent Sifre, Michal Valko, Simon Osindero, Timothy Lillicrap, Nicolas Heess, and Charles Blundell. Retrieval augmented reinforcement learning. In *International Conference on Machine Learning*, 2022.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang. Realm: Retrieval-augmented language model pre-training. In *International Conference on Machine Learning*, 2020.
- David Ha and Jurgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018.
- Danijar Hafner, Timothy Lillicrap, Ian Fischer, Ruben Villegas, David Ha, Honglak Lee, and James Davidson. Learning latent dynamics for planning from pixels. In *International Conference on Machine Learning*, 2019.
- Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. Dream to control: Learning behaviors by latent imagination. In *International Conference on Learning Representations*, 2020.
- Michael Janner, Justin Fu, Marvin Zhang, and Sergey Levine. When to trust your model: Model-based policy optimization. In *Advances in Neural Information Processing Systems*, 2019.
- Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. Generalization through memorization: Nearest neighbor language models. In *International Conference on Learning Representations*, 2020.
- Rahul Kidambi, Aravind Rajeswaran, Praneeth Netrapalli, and Thorsten Joachims. Morel: Model-based offline reinforcement learning. In *Advances in Neural Information Processing Systems*, 2020.

- Hung Le, Thommen Karimpanal George, Majid Abdolshah, Truyen Tran, and Svetha Venkatesh. Model-based episodic memory induces dynamic hybrid controls. In *Advances in Neural Information Processing Systems*, 2021.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen-tau Yih, Tim Rocktaschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*, 2020.
- Alexander Pritzel, Benigno Uria, Sriram Srinivasan, Adrià Puigdomènech Badia, Oriol Vinyals, Demis Hassabis, Daan Wierstra, and Charles Blundell. Neural episodic control. In *International Conference on Machine Learning*, 2017.
- Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Mastering atari, go, chess and shogi by planning with a learned model. *Nature*, 2020.
- Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling llm test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*, 2024.
- Richard S. Sutton. Integrated architectures for learning, planning, and reacting based on approximating dynamic programming. *Machine Learning*, 1990.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations*, 2023.
- Greg Wayne, Chia-Chun Hung, David Amos, Mehdi Mirza, Arun Ahuja, Agnieszka Grabska-Barwinska, Jack Rae, Piotr Mirowski, Joel Leibo, Adam Santoro, Mevlana Gemici, Malcolm Reynolds, Tim Harley, Josh Abramson, Shakir Mohamed, Danilo Rezende, David Saxton, Adam Cain, Chloe Hillier, David Silver, Koray Kavukcuoglu, Matthew Botvinick, Demis Hassabis, and Timothy Lillicrap. Unsupervised predictive memory in a goal-directed agent. In *International Conference on Machine Learning*, 2018.
- Yuhuai Wu, Markus N. Rabe, DeLesley Hutchins, and Christian Szegedy. Memorizing transformers. In *International Conference on Learning Representations*, 2022.
- Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. In *Advances in Neural Information Processing Systems*, 2023.
- Tianhe Yu, Garrett Thomas, Lantao Yu, Stefano Ermon, James Zou, Sergey Levine, Chelsea Finn, and Tengyu Ma. Mopo: Model-based offline policy optimization. In *Advances in Neural Information Processing Systems*, 2020.

A ADDITIONAL DIAGNOSTIC DEFINITION

The hallucinated-rollout indicator is one when the selected candidate’s model proxy is near-goal while its true return is poor. In code, this corresponds to a proxy score above -1.75 and true return below -5.0 . The threshold is used only for diagnostics; selection itself uses the raw proxy or repaired proxy.

B ANONYMOUS SUBMISSION NOTE

The source uses `iclr2026_conference.sty` and keeps `\iclrfinalcopy` commented. The resulting PDF displays the standard under-review anonymous author block.